

# Transformers

CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

# Questions

- How to summarize what it means for an ML system to be **trusted**?
- Name one of the four major characteristics of a **trustworthy** ML system according to the book's opinion.

Explain what does it mean and why it matters

# Questions

- Have you noticed any news, articles, policies, events that have implications in trustworthy deep learning in recent years?

# Readings

- Readings are updated on the website (syllabus page)
- **Readings:**
  - [Attention is All You Need](#)
  - [Generative pre-training](#)
  - [GNN Survey](#)
- This lecture is not explicitly tested
  - But in future lectures we will assume knowledge of this when developing trustworthy components on top of Transformers

# Outline of Today's Lecture

## **1. Self-Attention and Transformers**

## 2. Transformers Applications

## 3. Graph Neural Networks

# Sequence Learning

- Inputs from different domains can be seen as the general **sequence** of **tokens**

## Domain

## Sequence

## Token

## Structure

NLP


**Sentence:** [SOS, “Deep”,  
“learning”, “can”, “empower”,  
“sciences”, EOS]

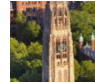
**Word:** “learning”  
**Phrase:** [“Deep”,  
“learning”]

Sequential correlations

CV

**Image:**






**Pixel**  
**Patch:** 

Spatial correlations

Graph

**Graph:**



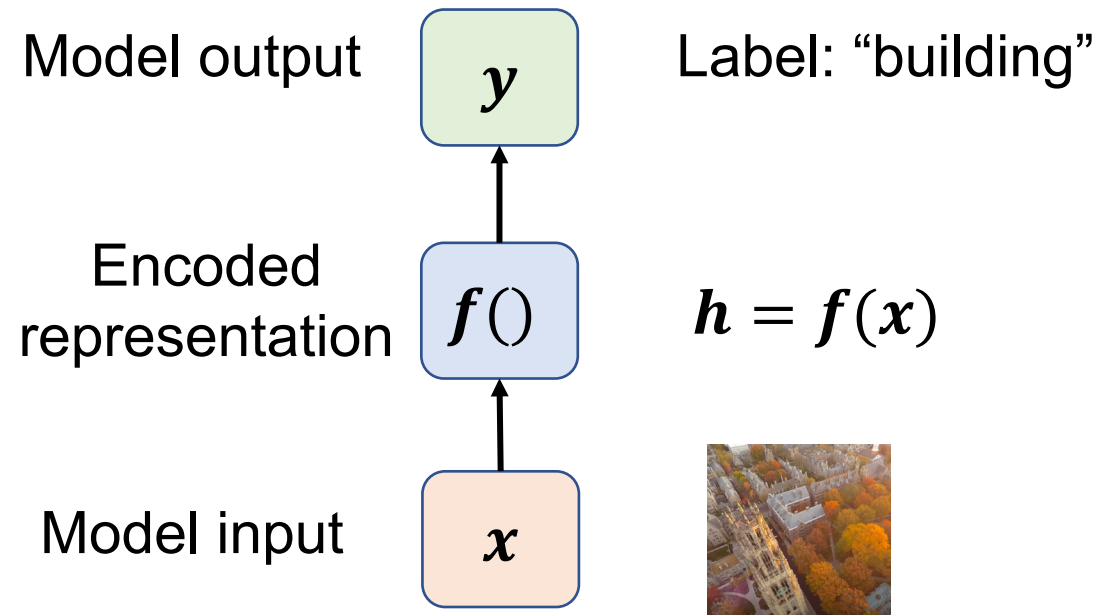
**Node:**   
**Subgraph:** 

Adjacency

# Standard Supervised Learning Setting

- **One (token) to One (token)**

- Input is a single token (e.g., an image), and the output is its attribute (e.g., label) or another token.
- $h = f(x)$ ,  $f()$  is the model to learn.



# Sequence Learning — Application

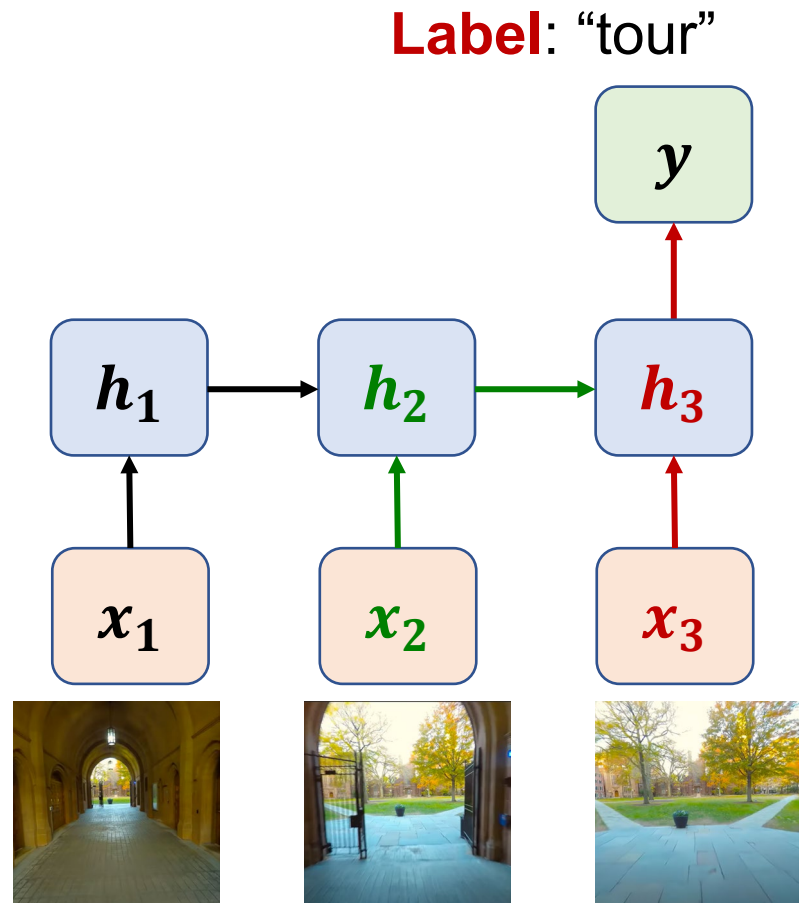
- **Many (tokens) to One**

- Input is a sequence of tokens (e.g. a video with frames), and the output is its attribute (e.g. label) or another token.
- $h_1 = f(x_1)$
- To generate  $h_2$ , we would like to incorporate both  $x_2$  and the preceding frame  $x_1$  and  $h_2 = f(x_2, h_1)$ . Here  $f()$  is shared across all timesteps

- $h_i = f(x_i, h_{i-1})$

Current token

Previous token






# Sequence Learning — Application

- We can process a sequence of tokens  $X = [x_1, x_2, \dots, x_n]$  by applying a recurrence formula at every time step
- Recurrent neural networks

$$h_i = f_W(x_i, h_{i-1})$$

↓ new state      ↓ current input      ↘ old state

- For example,  $h_i = \sigma(W_x x_i + W_h h_{i-1} + b_h)$ ,  
and  $y_i = \sigma(W_y h_i + b_y)$



A folded diagram of RNNs

# Sequence Learning — Application

- **Many (tokens) to Many**

- The sequence is first encoded into a hidden representation, then gradually decoded by the decoder.

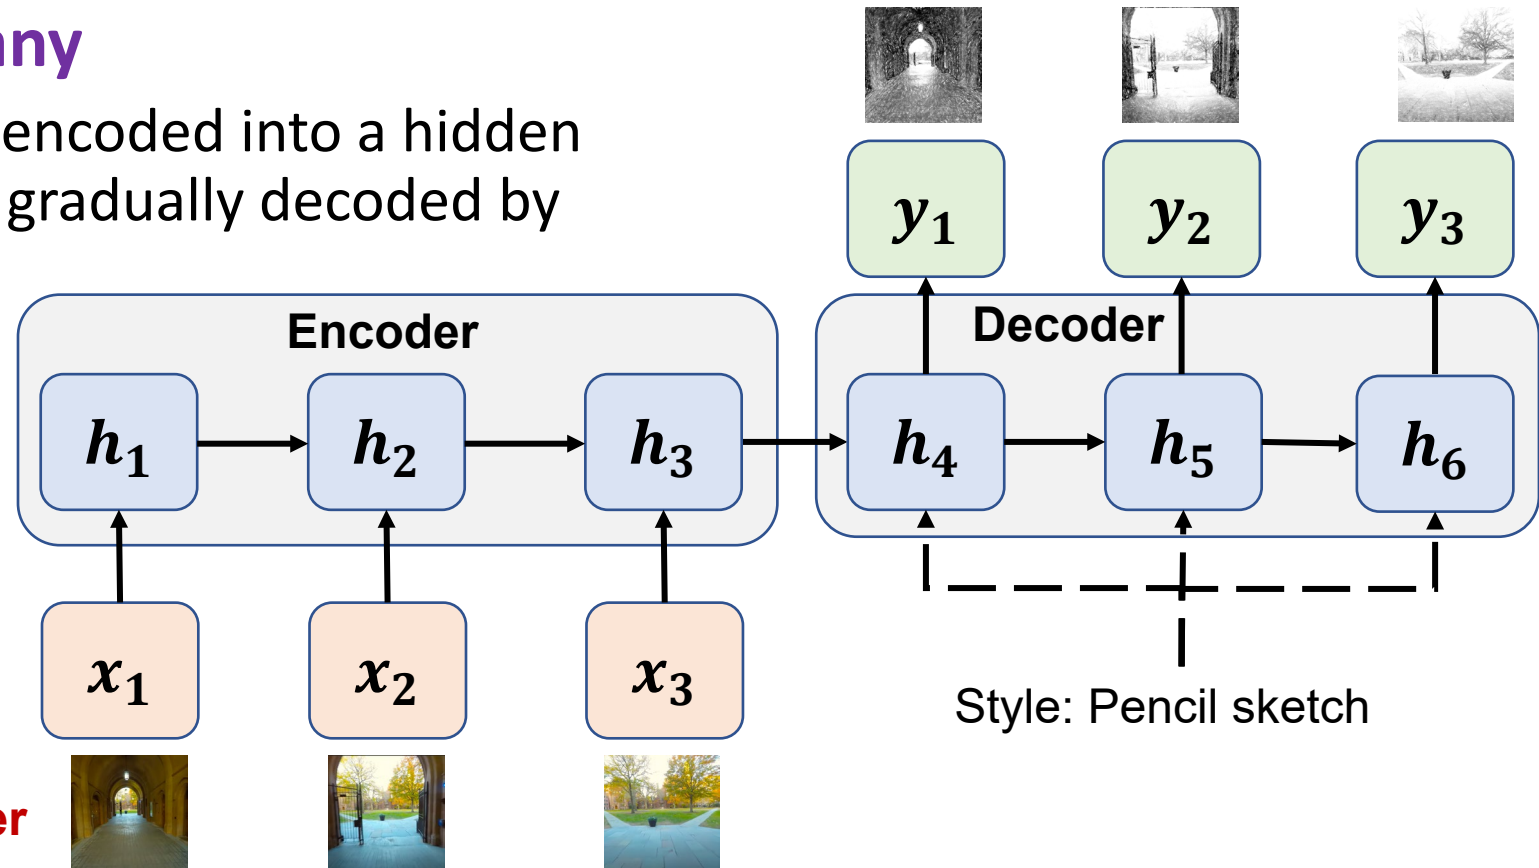


Diagram of video **style transfer**

# Sequence Learning — Application

- **One (token) to One (token)**

- **Many to One**


- Protein to property
- Sentence to sentiment
- Song to name

- **One to Many**


- Image to caption (multiple words)

- **Many to many**

- Translation: English to French
- Time series: history to future
- Graph autoregression




**Speech classification**  
(many-to-one)



**Image Captioning**  
(one-to-many)


# Sequence Learning



$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

**What are the issues and challenges of RNNs?**

# Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

## Problems of RNNs

- Sequential computation prevents parallelization
- Capacity of handling long sequences
- Mainly focusing on modeling recurrence
  - does not capture other correlations (hierarchical, long-range, polysemy.... ) well

# Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$

## Problems of RNNs


1. parallelization — — — — — ➤
2. long sequences — — — — — ➤
3. only recurrence — — — — — ➤

## Solutions by Transformers

1. **Parallel input:** Input All tokens at the same time
2. **Self-Attention:** Enable attention in long-range
3. **Positional Embeddings  $p_i$ :** Model all possible correlations




# Transformers — Overview



- **Original paper:** Attention is all you need [Vaswani et al., 2017].
- **Key component:** Multi-head self-attention
- **Other components** of a transformer layer: layer normalization, skip connection, position-wise feed-forward layer (FFN, or MLP)
- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
  - **Encoder-only** (e.g BERT)
  - **Encoder-Decoder** (e.g [BART](#))
  - **Decoder-only** (e.g GPT)
  - We will show an example later


# Transformers — Overview



- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
  - **Encoder-only** (e.g BERT)
    - Many-to-one classification / regression
    - Sentiment classification, document classification ...
    - Word / Sentence embeddings for downstream tasks (e.g. recommender system)
  - **Encoder-Decoder** (e.g [BART](#))
  - **Decoder-only** (e.g GPT)
  - We will show an example later



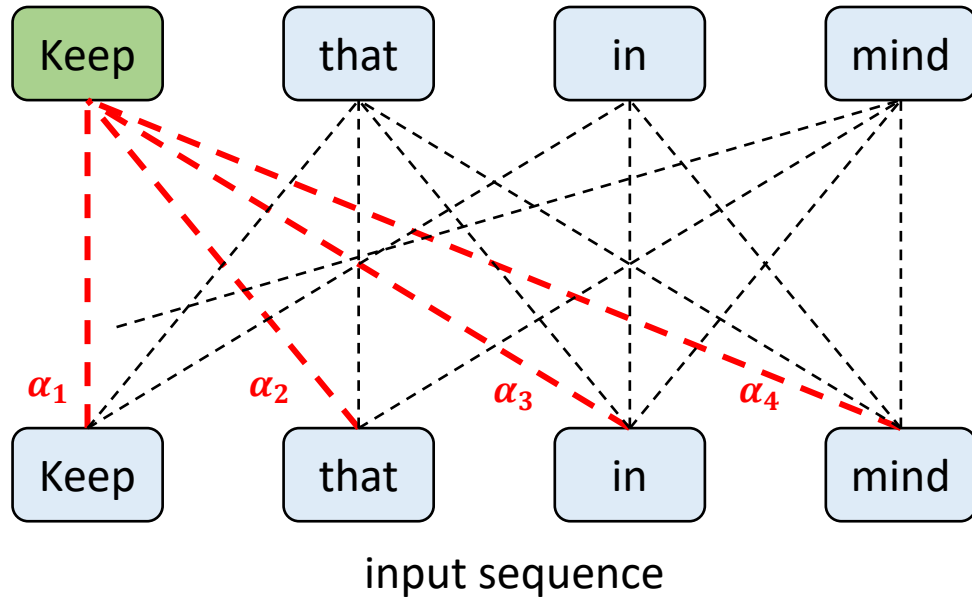
# Transformers — Overview



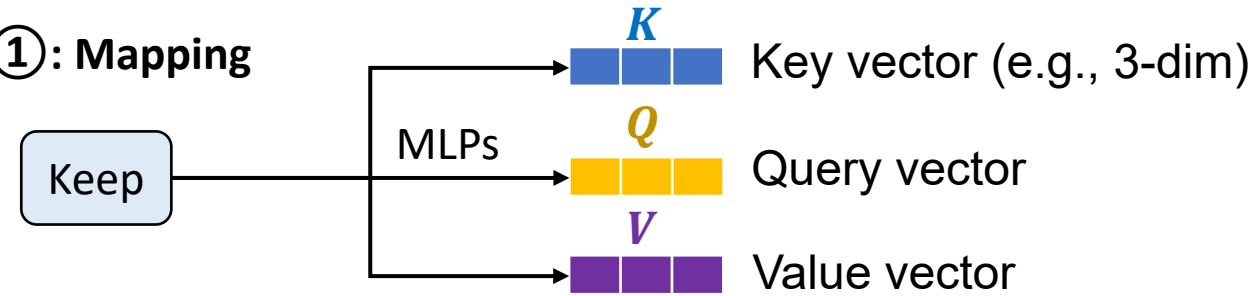
- **Model usage:** Pre-training followed by fine-tuning. The transferred model can be:
  - **Encoder-only** (e.g BERT)
  - **Encoder-Decoder** (e.g [BART](#))
    - Many-to-many use cases
    - Summarization, translation, style transfer ...
  - **Decoder-only** (e.g OpenAI GPT)
    - One-to-many use cases
    - Image / text / code generation, dialogue systems ...
    - GPT-3/4 based [apps](#)

# Transformers — Self-Attention (1/5)

## Example:



### Step ①: Mapping



### Step ②: Attention

$$\alpha_1, \alpha_2, \alpha_3, \alpha_4 = \text{Softmax} \left( \frac{Q_{\text{Keep}} \times \begin{bmatrix} K_{\text{Keep}} \\ K_{\text{that}} \\ K_{\text{in}} \\ K_{\text{mind}} \end{bmatrix}}{\text{Keep}} \right)$$



Keep that in mind

### Step ③: Update

$$V'_{\text{Keep}} = \alpha_1 \times V_{\text{Keep}} + \alpha_2 \times V_{\text{that}} + \alpha_3 \times V_{\text{in}} + \alpha_4 \times V_{\text{mind}}$$

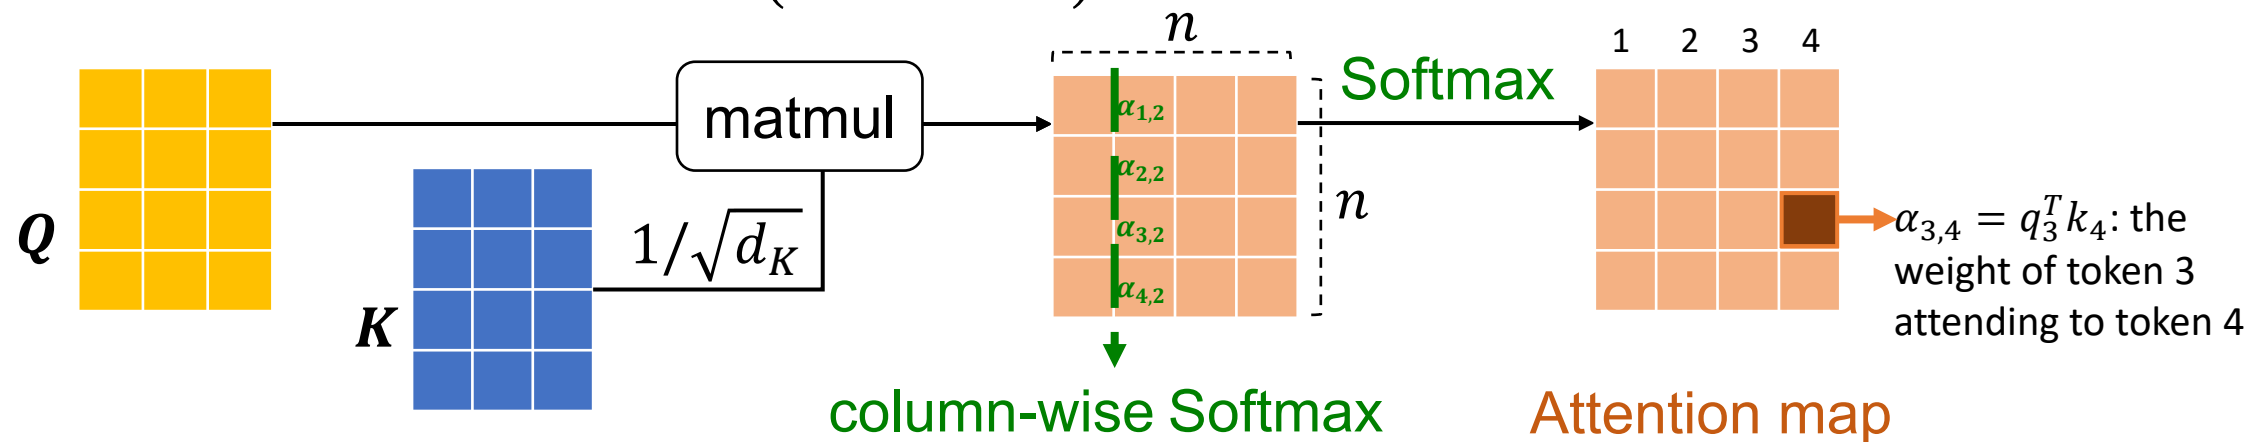
# Transformers — Self-Attention (2/5)

- Formally, given an input sequence  $X = [x_1, x_2, \dots, x_n] \in \mathbb{R}^{n \times d_X}$
- Step ①: Query  $Q = XW_Q$ , Key  $K = XW_K$ , Value  $V = XW_V$ 
  - $W_K \in \mathbb{R}^{d_X \times d_K}$ , and thus  $K \in \mathbb{R}^{n \times d_K}$
  - We require  $d_K = d_Q$ , for simplicity, we set  $d_K = d_Q = d_V := d$



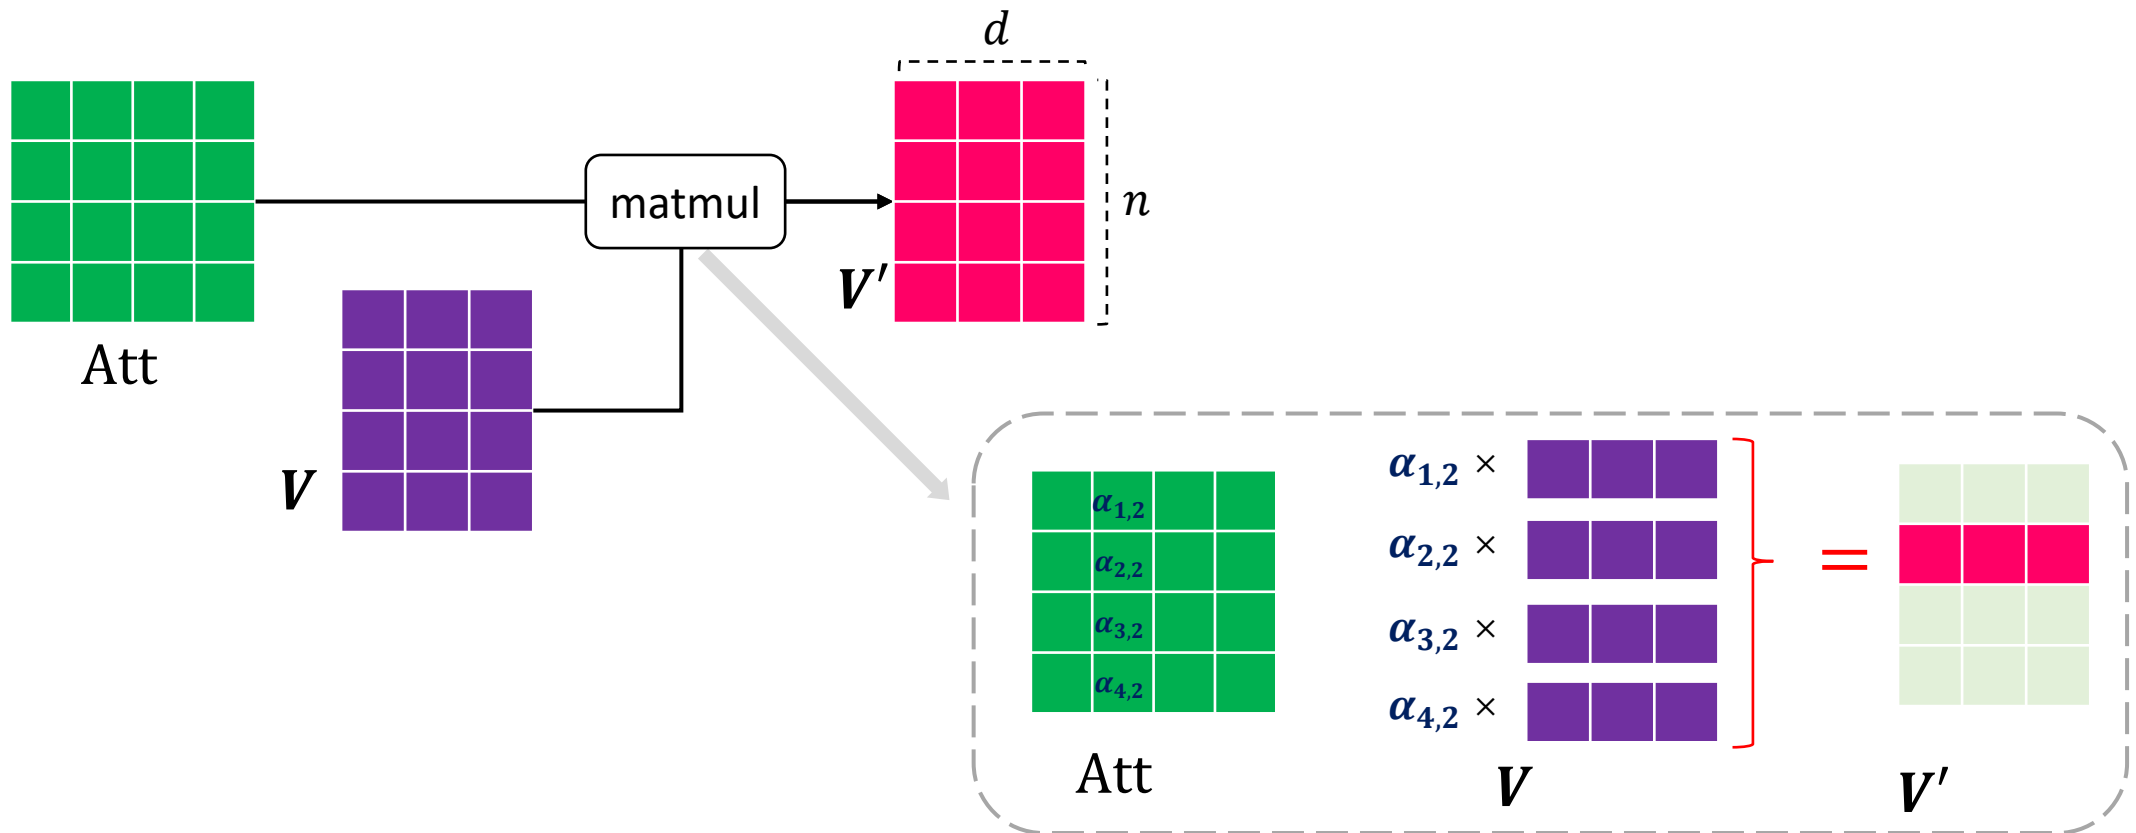
# Transformers — Self-Attention (3/5)

- Step ② : Attention map  $\text{Att} = \text{Softmax}\left(\frac{QK^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$  (Softmax is col-wise)
  - The matrix multiplication  $QK^T$  performs dot-product for every possible pair of queries and keys, resulting in an attention map.
  - **Normalization factor**  $1/\sqrt{d_K}$  : performing dot-product over two vectors with variance  $\sigma^2$  results in a scalar having  $d_K$ -times higher variance,
    - $q \sim N(0, \sigma^2), k \sim N(0, \sigma^2) \rightarrow \text{Var}\left(\sum_{i=1}^{d_K} q[i]k[i]\right) = \sigma^4 d_K$




# Transformers — Self-Attention (4/5)

- Step ③: Updated value  $V' = \text{Att } V \in \mathbb{R}^{n \times d}$  Matrix product



# Transformers — Self-Attention (5/5)

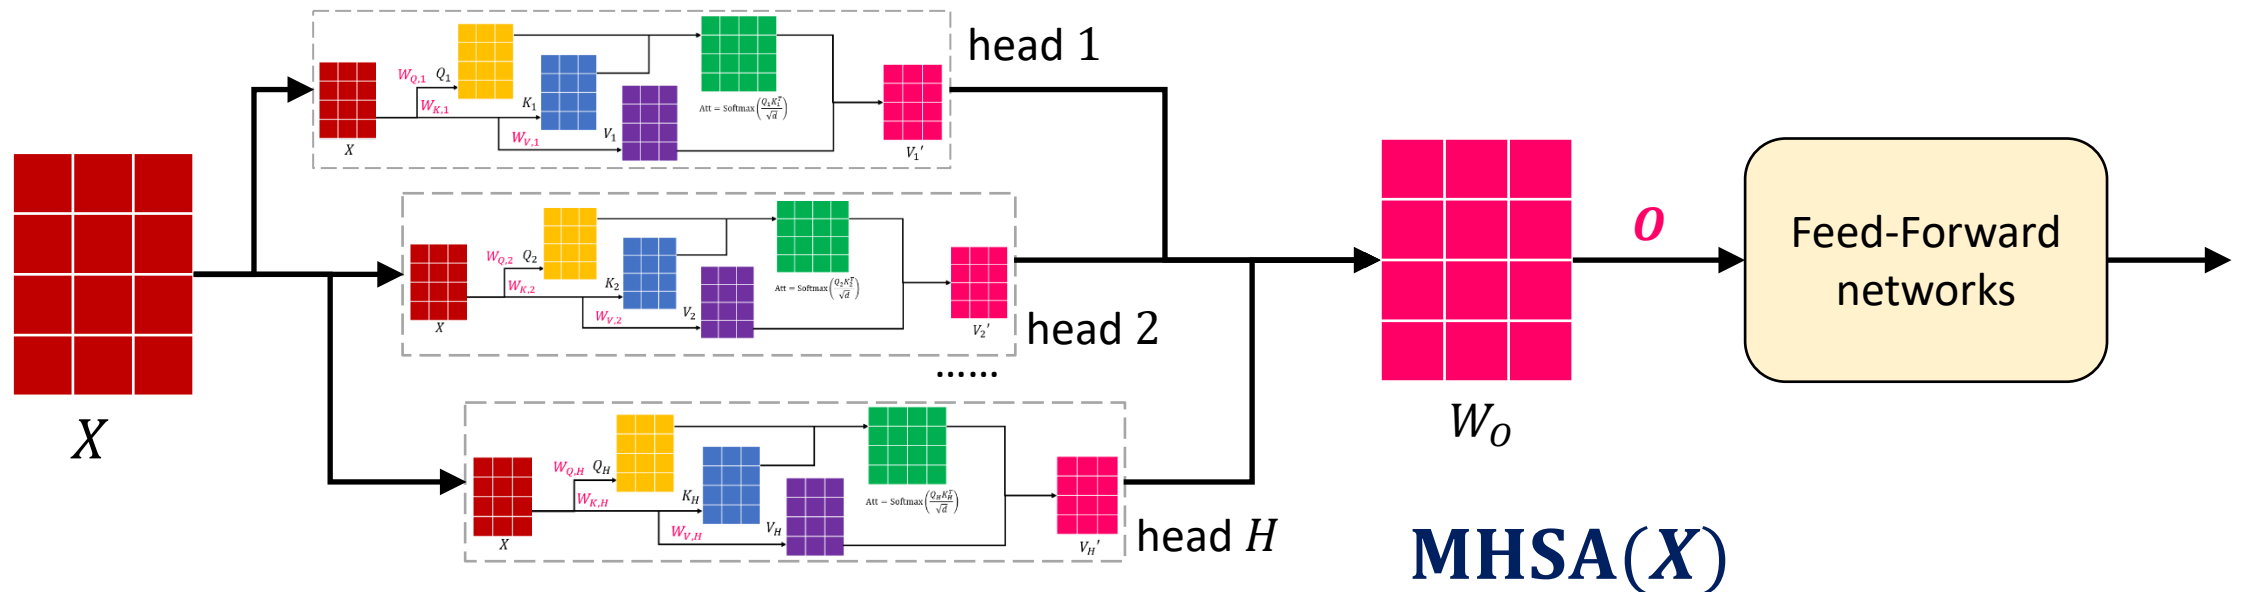
- Putting all together



**The computation complexity is quadratic to number of tokens**


# Transformers — Multi-Head Self-Attention

- There are usually **multiple aspects** that a token can attend to.
- We extend the attention to multiple heads, with multiple  $(Q, K, V)$  triplets on the same features.
  - The output of multi-head self-attention  $O = \text{Concat}([V'_1, V'_2, \dots, V'_H])W_O$
  - Learnable parameters in each attention layer:  $W_{Q,i}, W_{K,i}, W_{V,i} \in R^{d_x \times d}$  for head  $i$ ,  $W_O \in R^{Hd \times d_o}$



# Transformers — Layer (1)

- **MHSA**: multi-head self-attention
- Transformer layer:  $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X))$
- **Residual connections** are added to
  - Enable smooth gradient flow in deep transformers
  - Keep the information of the original sequence.




**What are some advantages or challenges when trying to explain Transformer-based models?**



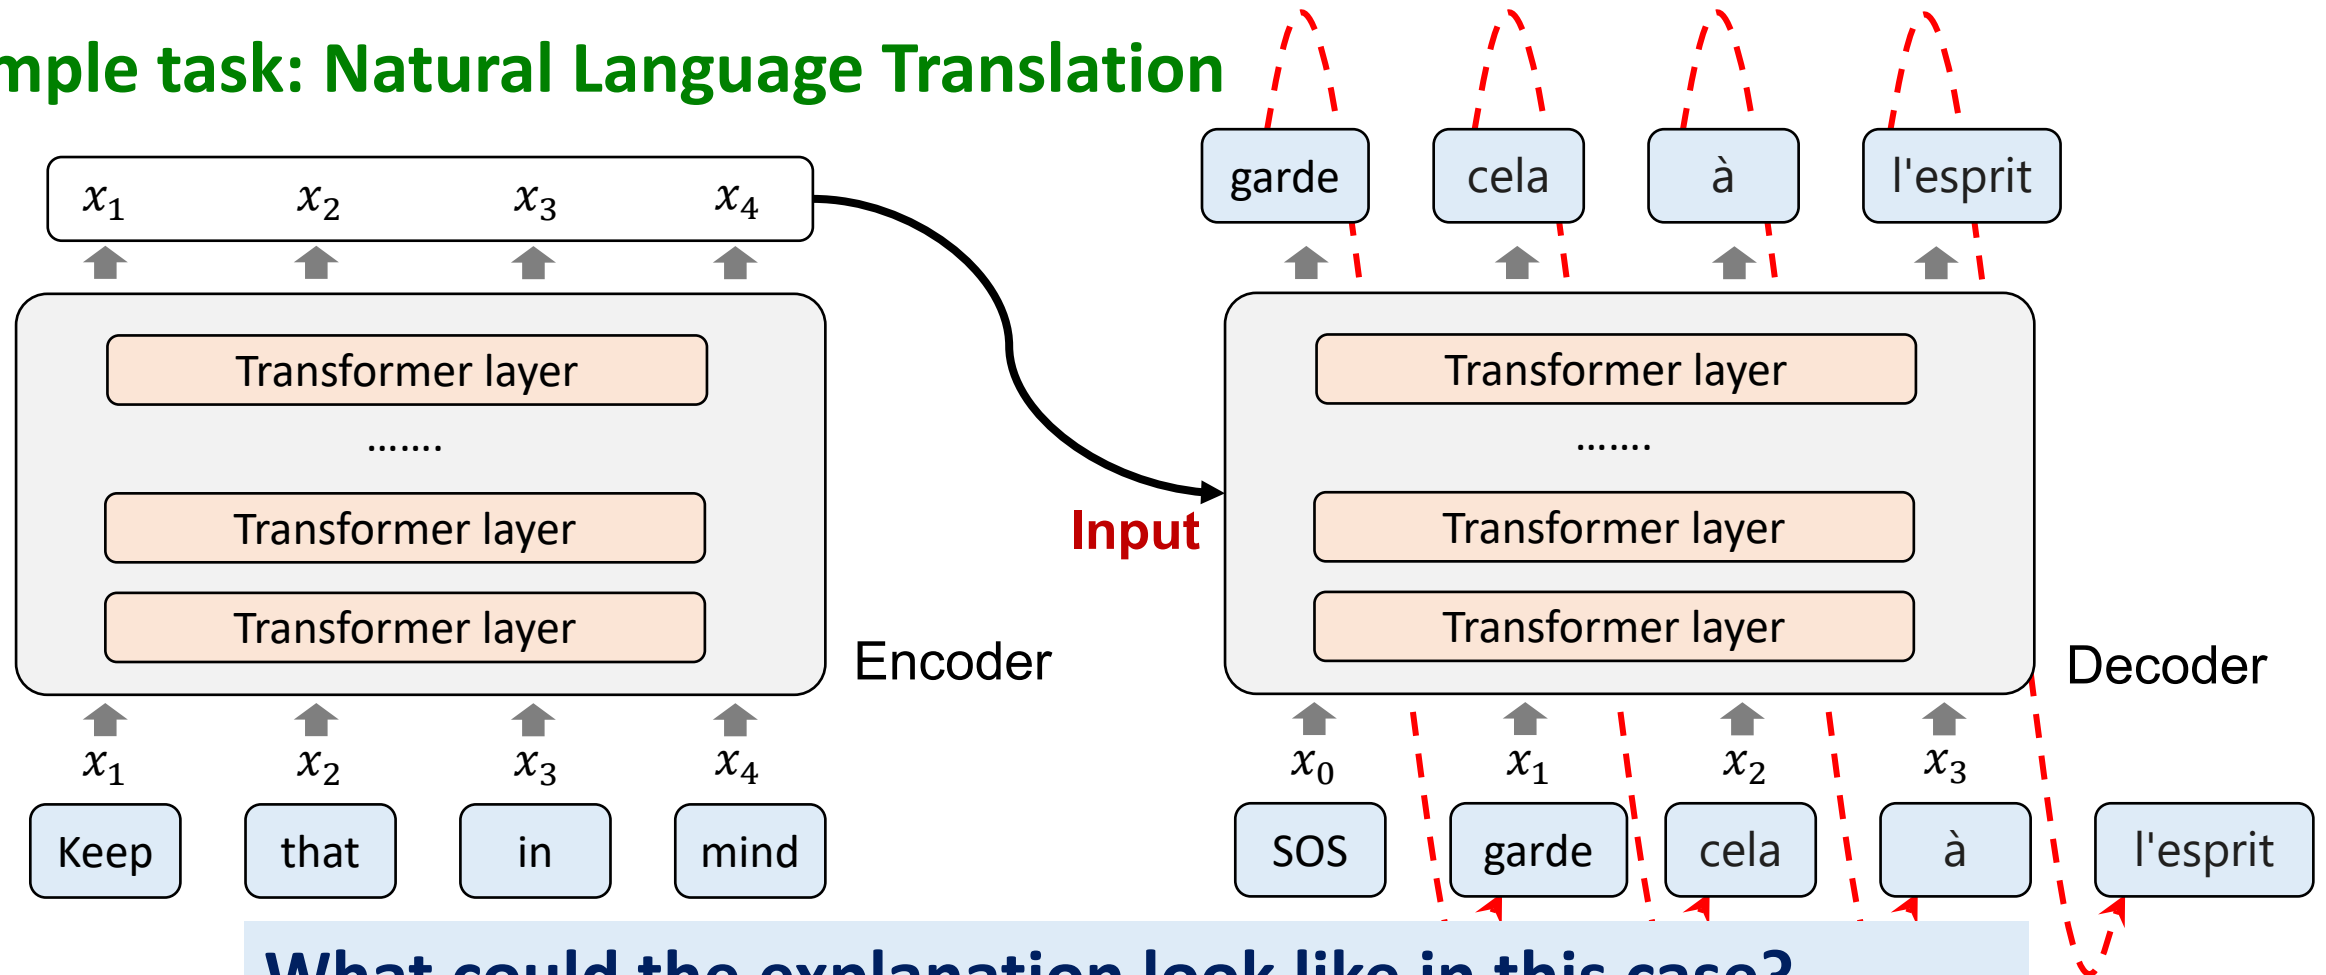
# Transformers — Layer (2)

- Transformer layer:  $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X)) \rightarrow \text{LayerNorm}(X + \text{FFN}(X))$
- Layer Normalization** is used to enable faster training with small regularization and keep features in similar magnitudes.
  - BatchNorm isn't applied because batch size is usually small in Transformers due to GPU memory constraints. Besides, BatchNorm has been shown to lead to worse performance in NLP.
- MLPs** are added for “post-processing”, and allow transformations on each sequence token.



# Transformers —Encoder / Decoder

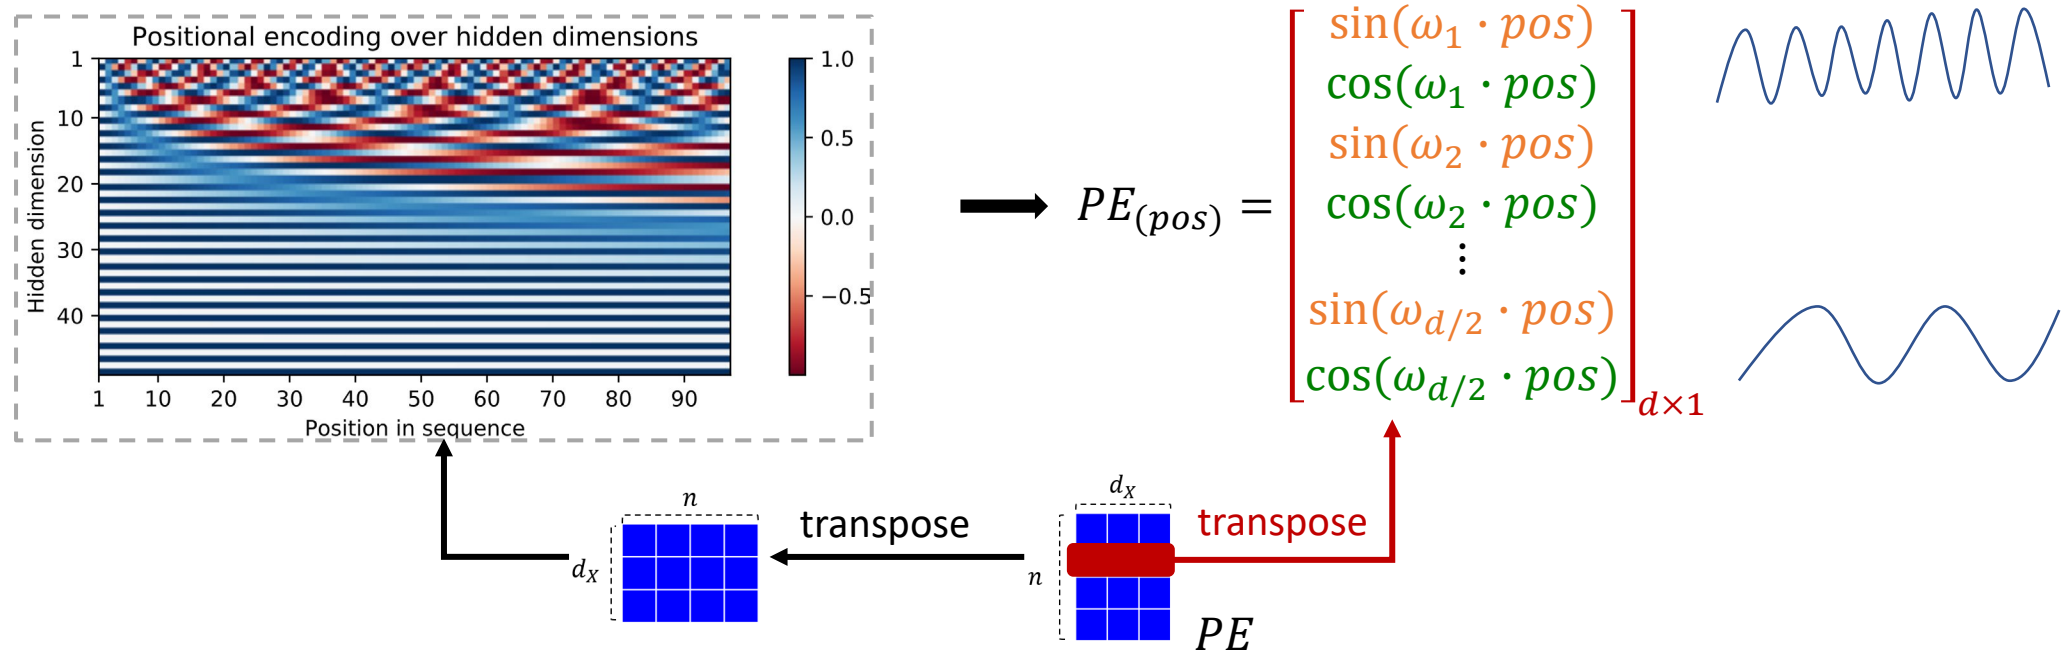
## Example task: Natural Language Translation



# Transformers —Positional Encoding (3)

- **Cosine encoding**

- $PE_{(pos,2i)} = \sin(pos/10000^{2i/d_x})$ ,  $PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_x})$ .
- $\omega_i = 1/10000^{2i/d_x}$ .
- Relative distance:  $PE_{(pos+k)}$  can be easily represented as a linear function of  $PE_{(pos)}$  (show it).



# Summary: Transformer Architecture

- Multi-Head Self-Attention (**MHSA**( $X$ ))
  - For head  $i$ 
    - $Q_i = XW_{Q_i}, K_i = XW_{K_i}, V_i = XW_{V_i}$
    - $\text{Att}_i = \text{Softmax}\left(\frac{Q_i K_i^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$
    - $V_i' = \text{Att}_i V_i \in \mathbb{R}^{n \times d}$
  - Concatenating all heads:  $O = \text{Concat}([V_1', V_2', \dots, V_H'])W_O$
- $X = \text{LayerNorm}(X + \text{MHSA}(X))$
- $X = \text{LayerNorm}(X + \text{FFN}(X))$

# Outline of Today's Lecture

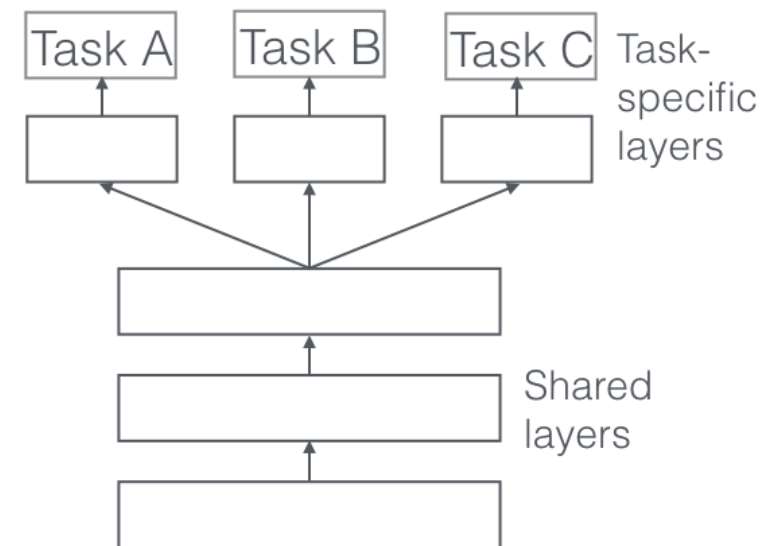
1. Self-Attention and Transformers

**2. Transformers Applications**

3. Graph Transformers and Sparse Transformers

# Why is Transformer a Popular Choice

- Resolves various challenges of RNN-based architectures
- Attention makes the architecture **expressive and flexible** for different application scenarios
- It is very amenable to **self-supervised objectives**
  - We can leverage the vast number of **unsupervised examples** to learn a general model
  - Can be fine-tuned for **many downstream tasks**
  - Can out-perform models that are only trained for a specific downstream task



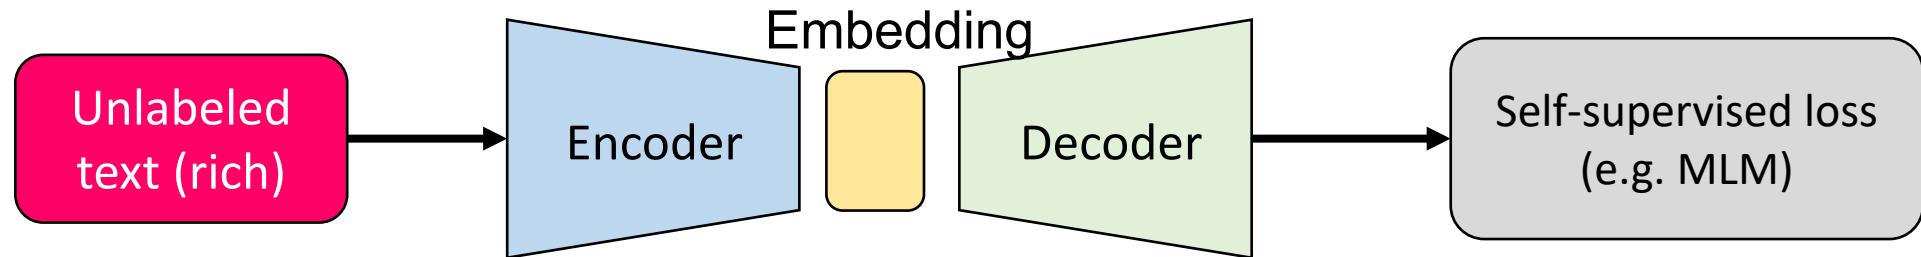
# Label Scarcity

- ML models are hungry to data, especially labeled data for supervised task.
- The fast development of computer vision largely benefits from **ImageNet**. It contains 14 million images **hand-annotated** by a team of researchers.
- This is often not possible for many domains. Most of time, it's easy to collect rich unlabeled data, but hard to obtain labeled data.
- **Solution:** **Pre-training** general-purpose language model on unlabeled large corpora (billions of characters) in **self-supervised** setting, then **fine-tuning** on smaller-scale tasks.

# Pre-training and Fine-tuning

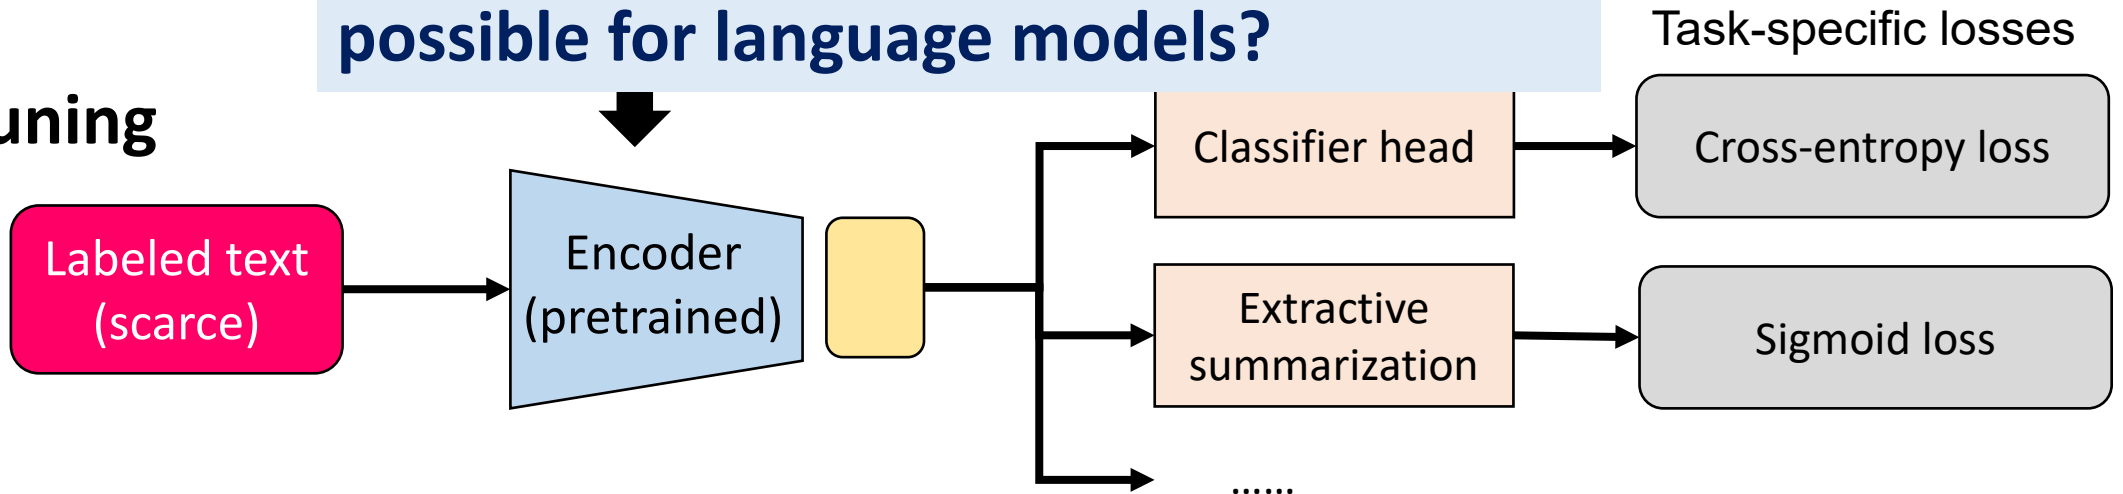
- **Pre-training**

## Encoder-only Pre-training Scenario



**What are some attacks that could be possible for language models?**

- **Fine-tuning**






# Transformers in NLP — BERT

BERT — **Bidirectional** Encoder Representations from Transformers [Devlin et al., 2018]

- **Pre-training task** (unsupervised): **Masked Language Model (MLM)**
  - First randomly masking  $m\%$  tokens in the input sequence.
    - In BERT, 15% tokens are masked at random (replaced with the special [MASK] token)
  - Predicting masked tokens using remaining tokens.
  - Two modes: **Unidirectional** and **Bidirectional**.

**Unidirectional** [Radford et al., 2018]



- Maximize Likelihood of “all” given “Att” and “is”

**Bidirectional**



- Maximize Likelihood of “all”, given “Att” , “is”, “you”, “need”.

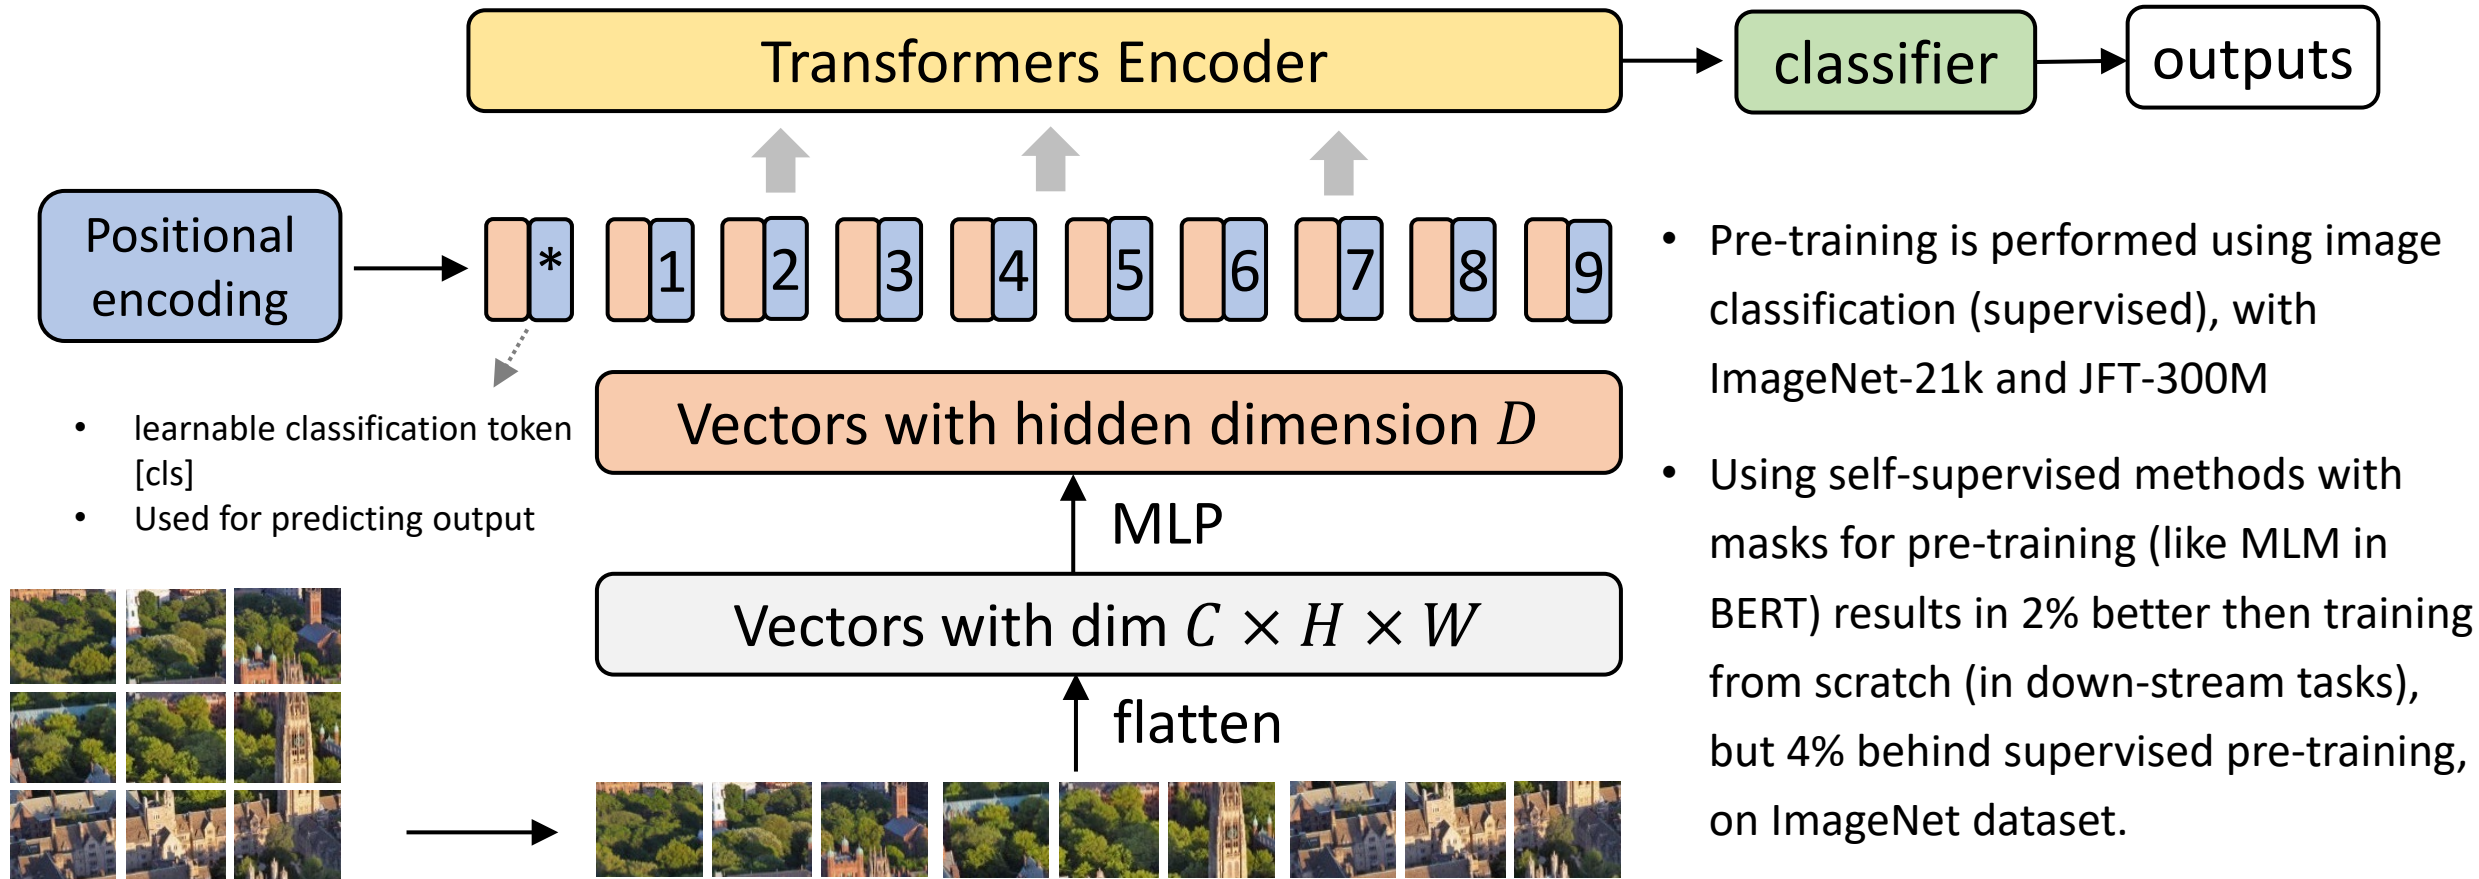
# Transformers in NLP — RoBERTa

**RoBERTa** — **Robustly** Optimized BERT [Liu et al., 2019]

- **Pretraining data**: BooksCorpus (800 M words) [Zhu et al., 2015], English Wiki (2500 M words), CC-News, OpenWebText [Gokaslan and Cohen, 2019], Stories [Trinh and Le, 2018]
  - Partition the corpus into “sentences” with fixed length of 512 tokens.
- **Hyperparameters** in use (also commonly adopted in most NLP Transformers):
  - **12-Layer** Encoder + **12-Layer** Decoder  
(Pretrained Encoder is used more frequently in down-stream tasks)
  - Hidden dimension **768** = 12 (num of Heads)  $\times$  64 (dim of Head)
  - Learning rate: Warmup then linear decay
    - Warmup: Gradually increasing the learning rate to a specific value in the first few epochs
    - Linear decay: Decreasing the learning rate by the same amount (decrement) every epoch.

# Transformers in CV — ViT [Dosovitskiy et al., ICLR 2021]

- An image patch is treated as a word in this context, and an image is partitioned to  $16 \times 16$  tokens.



- Pre-training is performed using image classification (supervised), with ImageNet-21k and JFT-300M
- Using self-supervised methods with masks for pre-training (like MLM in BERT) results in 2% better than training from scratch (in down-stream tasks), but 4% behind supervised pre-training, on ImageNet dataset.

# Outline of Today's Lecture

**1. Self-Attention and Transformers**

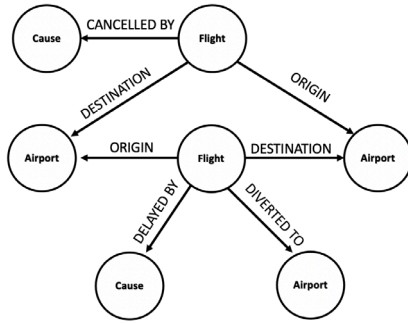
**2. Transformers Applications**

**3. Graph Neural Networks**

# Types of Networks and Graphs (1)

- **Networks (also known as Natural Graphs):**
  - **Social networks:**
    - **Society** is a collection of 7+ billion individuals
  - **Communication and transactions:**
    - Electronic devices, phone calls, financial transactions
  - **Biomedicine:**
    - Interactions between **genes/proteins** regulate life
  - **Brain connections:**
    - Our **thoughts** are hidden in the connections between billions of neurons

# Many Types of Data are Graphs (1)




**Event Graphs**



Image credit: [SalientNetworks](#)

**Computer Networks**



**Disease Pathways**




Image credit: [Wikipedia](#)

**Food Webs**



Image credit: [Pinterest](#)

**Particle Networks**

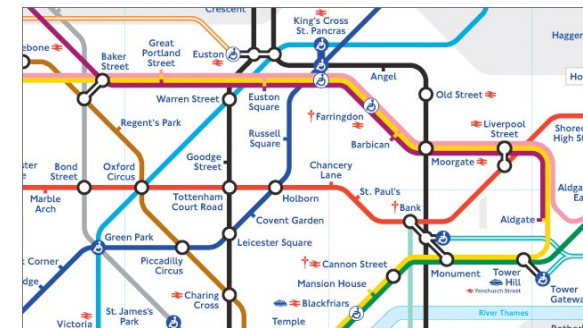


Image credit: [visitlondon.com](#)

**Underground Networks**



# Many Types of Data are Graph (2)



Image credit: [Medium](#)

**Social Networks**

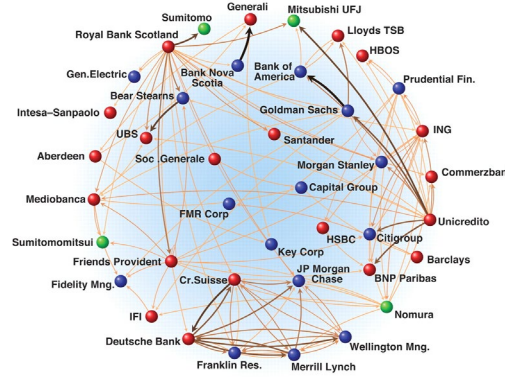



Image credit: [Science](#)

**Economic Networks**



Image credit: [Lumen Learning](#)

**Communication Networks**



**Citation Networks**



Image credit: [Missoula Current News](#)

**Internet**





Image credit: [The Conversation](#)

**Networks of Neurons**

# Graph Machine Learning Tasks: Overview


- Node-level prediction
- Link-level prediction
- Graph-level prediction
- Graph generation
  - Generative model





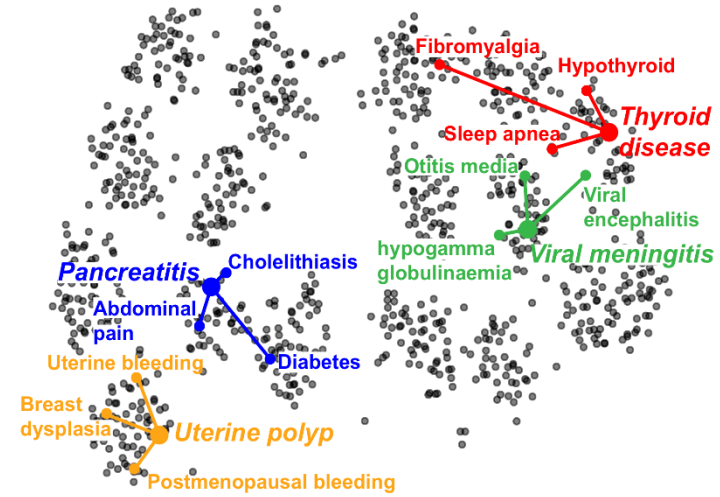
# Node Embeddings

- Intuition: Map nodes to d-dimensional **embeddings** (which are “**representations**” of nodes) such that similar nodes in the graph are embedded close together



Input graph

=




2D node embeddings

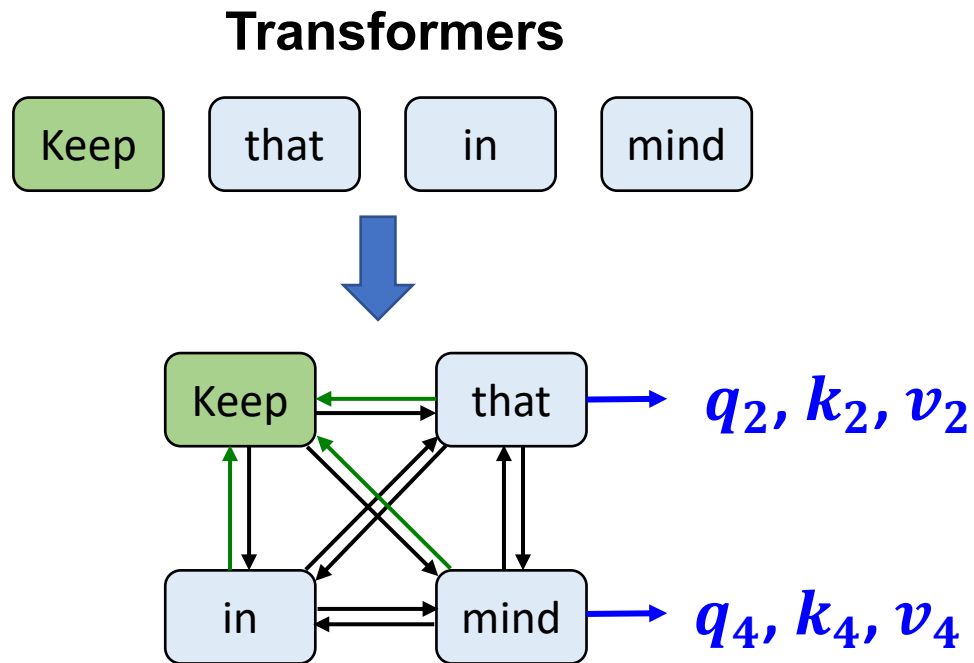
How to learn the mapping function  $f$ ?

# Deep Graph Encoders (1)

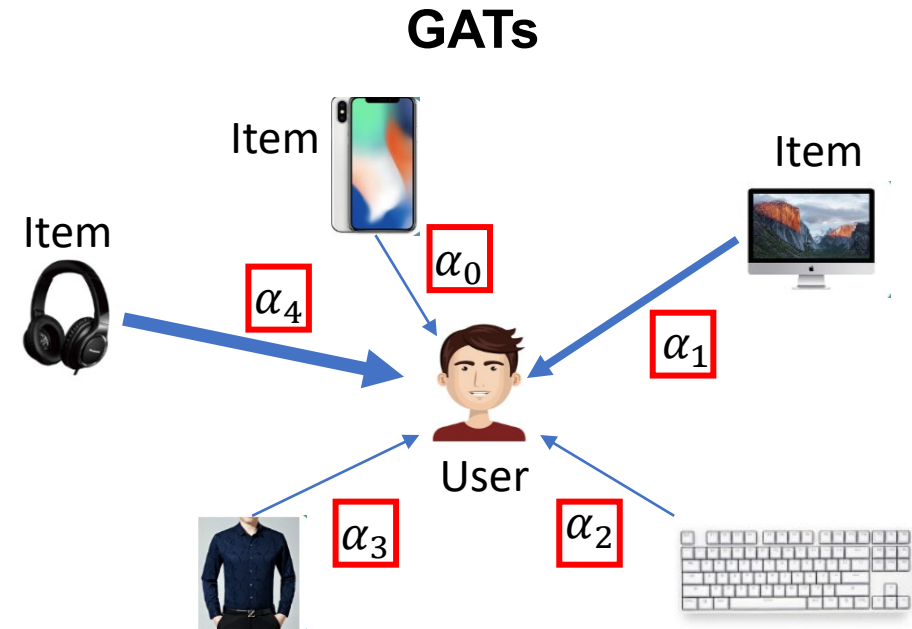
$\text{ENC}(\cdot)$  = multiple layers of  
non-linear transformations  
based on graph structures



# Transformers — in the Language of Graphs (1)



**Step ① Mapping:** Each node feature  $x_i$  is projected to  $q_i, k_i, v_i$ .




**Attention** computation: calculate the importance of neighbors

$$\alpha_{vu} = att \left( \mathbf{h}_v^{(l-1)}, \mathbf{h}_u^{(l-1)} \right)$$

# Transformers — in the Language of Graphs (2)


## Transformers



**Step ② Attention:** Calculate the edge weights using  $\mathbf{q}_i, \mathbf{k}_j$  of the two endpoints node  $i$  and  $j$  as  $e_{ij} = \mathbf{q}_i^T \mathbf{k}_j / \sqrt{d}$ , then normalizing it by the neighbors of node  $i$

$$\alpha_{ij} = \text{softmax}_i(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k \in N_i} \exp(e_{ik})}$$

## GATs




**Message** computing: transform information of neighbor node to a message


$$\mathbf{m}_u^{(l)} = \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_v^{(l-1)}, u \in N_v$$

# Transformers — in the Language of Graphs (3)

Transformers



GATs



**Step ③ Update:** Update each node feature according to its neighbors as

$$\mathbf{x}_i' = \sum_{k \in N_i} \alpha_{ik} \mathbf{x}_k$$

**Aggregate** message: aggregate messages from neighbor nodes

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N_v} \mathbf{m}_u^{(l)} \right)$$

# Transformers — in the Language of Graphs (4)

Summary: Comparison of **Self-attention (SA)** and **Graph Attention Networks (GAT)**

- Step ① Mapping
  - **SA**: different weights for  $q, k, v$ .  $q = w_q x, k = w_k x, v = w_v x$ .
  - **GAT**: shared weights for  $q, k, v$ .  $q = wx, k = wx, v = wx$ .
- Step ② Attention: **SA** uses dot-product attention, while (the original) **GAT** uses concatenation with MLP
  - Dot-product:  $e_{ij} = q_i^T k_j / \sqrt{d}$
  - Concat:  $e_{ij} = \text{act}(W [q_i || k_j])$ , where  $c$  is a weight vector and  $\text{act}$  is the activation function like LeakyReLU

# Graph Attention — in the Language of Transformer

- The above computations do not require the assumption of **the complete graph**.
  - We assume full connectivity, mostly because we do not want to miss any potential token correlations.
- Self-attention can be easily adapted to graph-structured input data where the token correlations are given by the **adjacency matrix**, by replacing the **complete graph** with the **input graph**.
  - $\text{Self-Att}(X) = \text{Softmax} \left( \frac{(\mathbf{W}_k X) (\mathbf{W}_q X)^T}{\sqrt{d}} \odot \mathbf{A}_G \odot \mathbf{W}_E \right) V$ .
  - $\mathbf{A}_G$  is the adjacency matrix of the graph and  $\mathbf{E}$  is the edge weights of the graph if any.
- The complexity is no longer  $O(n^2 d)$  but is linear to the edge number  $O(E)$

# Idea: Aggregate Neighbors (1)


- **Key idea:** Generate node embeddings based on **local network neighborhoods**





# Idea: Aggregate Neighbors (2)

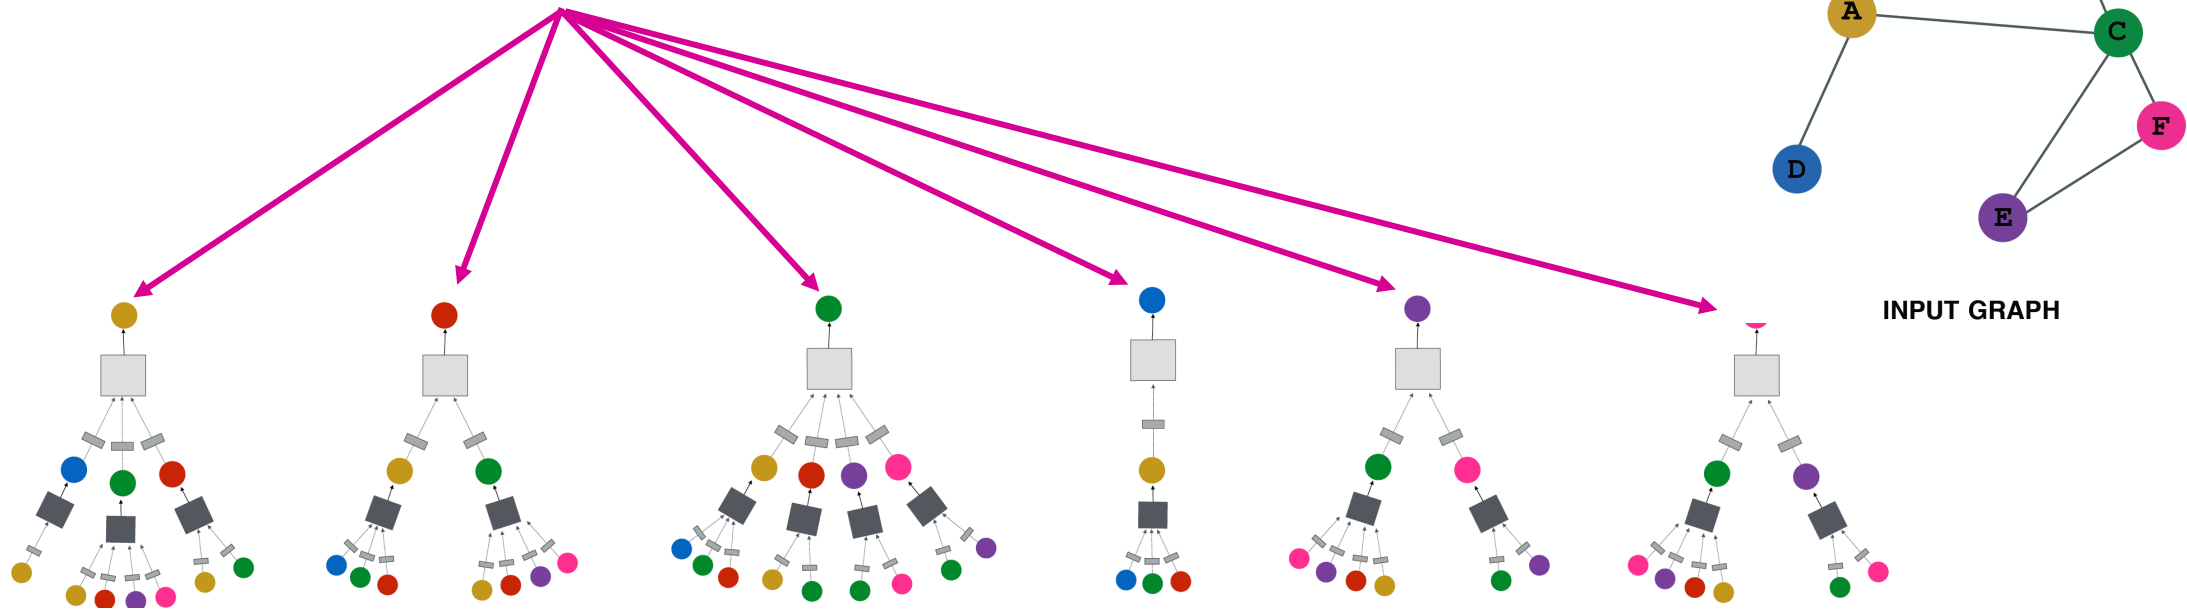
- **Intuition:** Nodes aggregate information from their neighbors using neural networks



# Idea: Aggregate Neighbors (3)


- **Intuition:** Network neighborhood defines a **computation graph**

Every node defines a computation graph based on its neighborhood!



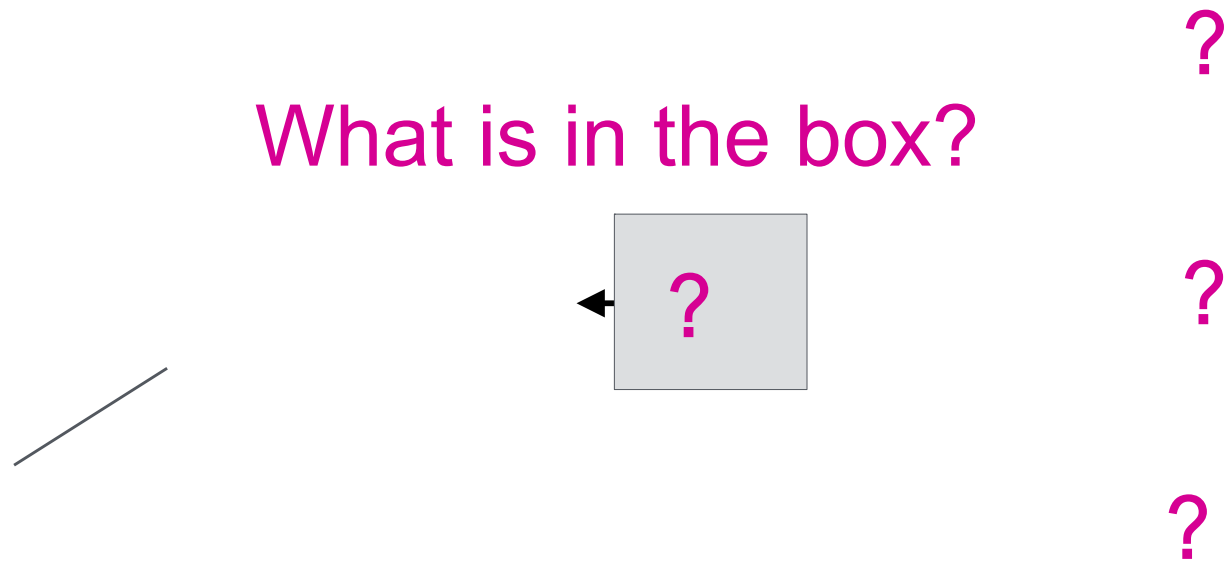
# Deep Model: Many Layers

- Model can be **of arbitrary depth**:
  - Nodes have embeddings at each layer
  - Layer-0 embedding of node  $u$  is its input feature,  $x_u$
  - Layer- $k$  embedding gets information from nodes that are  $k$  hops away



# Neighborhood Aggregation (1)


- **Neighborhood aggregation:** Key distinctions are in how different approaches aggregate information across the layers



# Neighborhood Aggregation (2)

- **Basic approach:** Average information from neighbors and apply a neural network

(1) average messages from neighbors



# Setup

- **Assume we have a graph  $G$ :**
  - $V$  is the **vertex set**
  - $A$  is the **adjacency matrix** (assume binary)
  - $X \in \mathbb{R}^{d \times |V|}$  is a matrix of **node features**
  - $v$ : a node in  $V$ ;  $N(v)$ : the set of neighbors of  $v$ .
  - **Node features:**
    - Social networks: User profile, User image
    - Biological networks: Gene expression profiles, gene functional information
    - When there is no node feature in the graph dataset:
      - Indicator vectors (one-hot encoding of a node)
      - Vector of constant 1:  $[1, 1, \dots, 1]$

# The Math: Deep Encoder

- **Basic approach:** Average neighbor messages and apply a neural network

